

# MLOps IRL: Vertex AI and Kubeflow Pipelines

Code breakfast

Julian de Ruiter & Timo Uelen



— DATA & AI CONSULTANCY AND TRAINING SERVICES

# Driving Your Success With Data and AI



# About us

- Julian de Ruiter

- Background in computer and life sciences
- Machine Learning Engineer at GoDataDriven since 2018
- Co-author of Manning's 'Data Pipelines with Apache Airflow'



Julian


- Timo Uelen

- Background in Information Sciences
- Machine Learning Engineer at GoDataDriven since 2020



Timo

# This morning



- 08:30 – 08:45 – Intro to MLOps + Vertex AI
- 08:45 – 09:00 – Getting started: Vertex Workbench
- 09:00 – 09:15 – Intro to Kubeflow pipelines
- 09:15 – 10:15 – Hackathon: building an ML pipeline
- 10:15 – 10:30 – Discussion & wrap-up

# MLOps

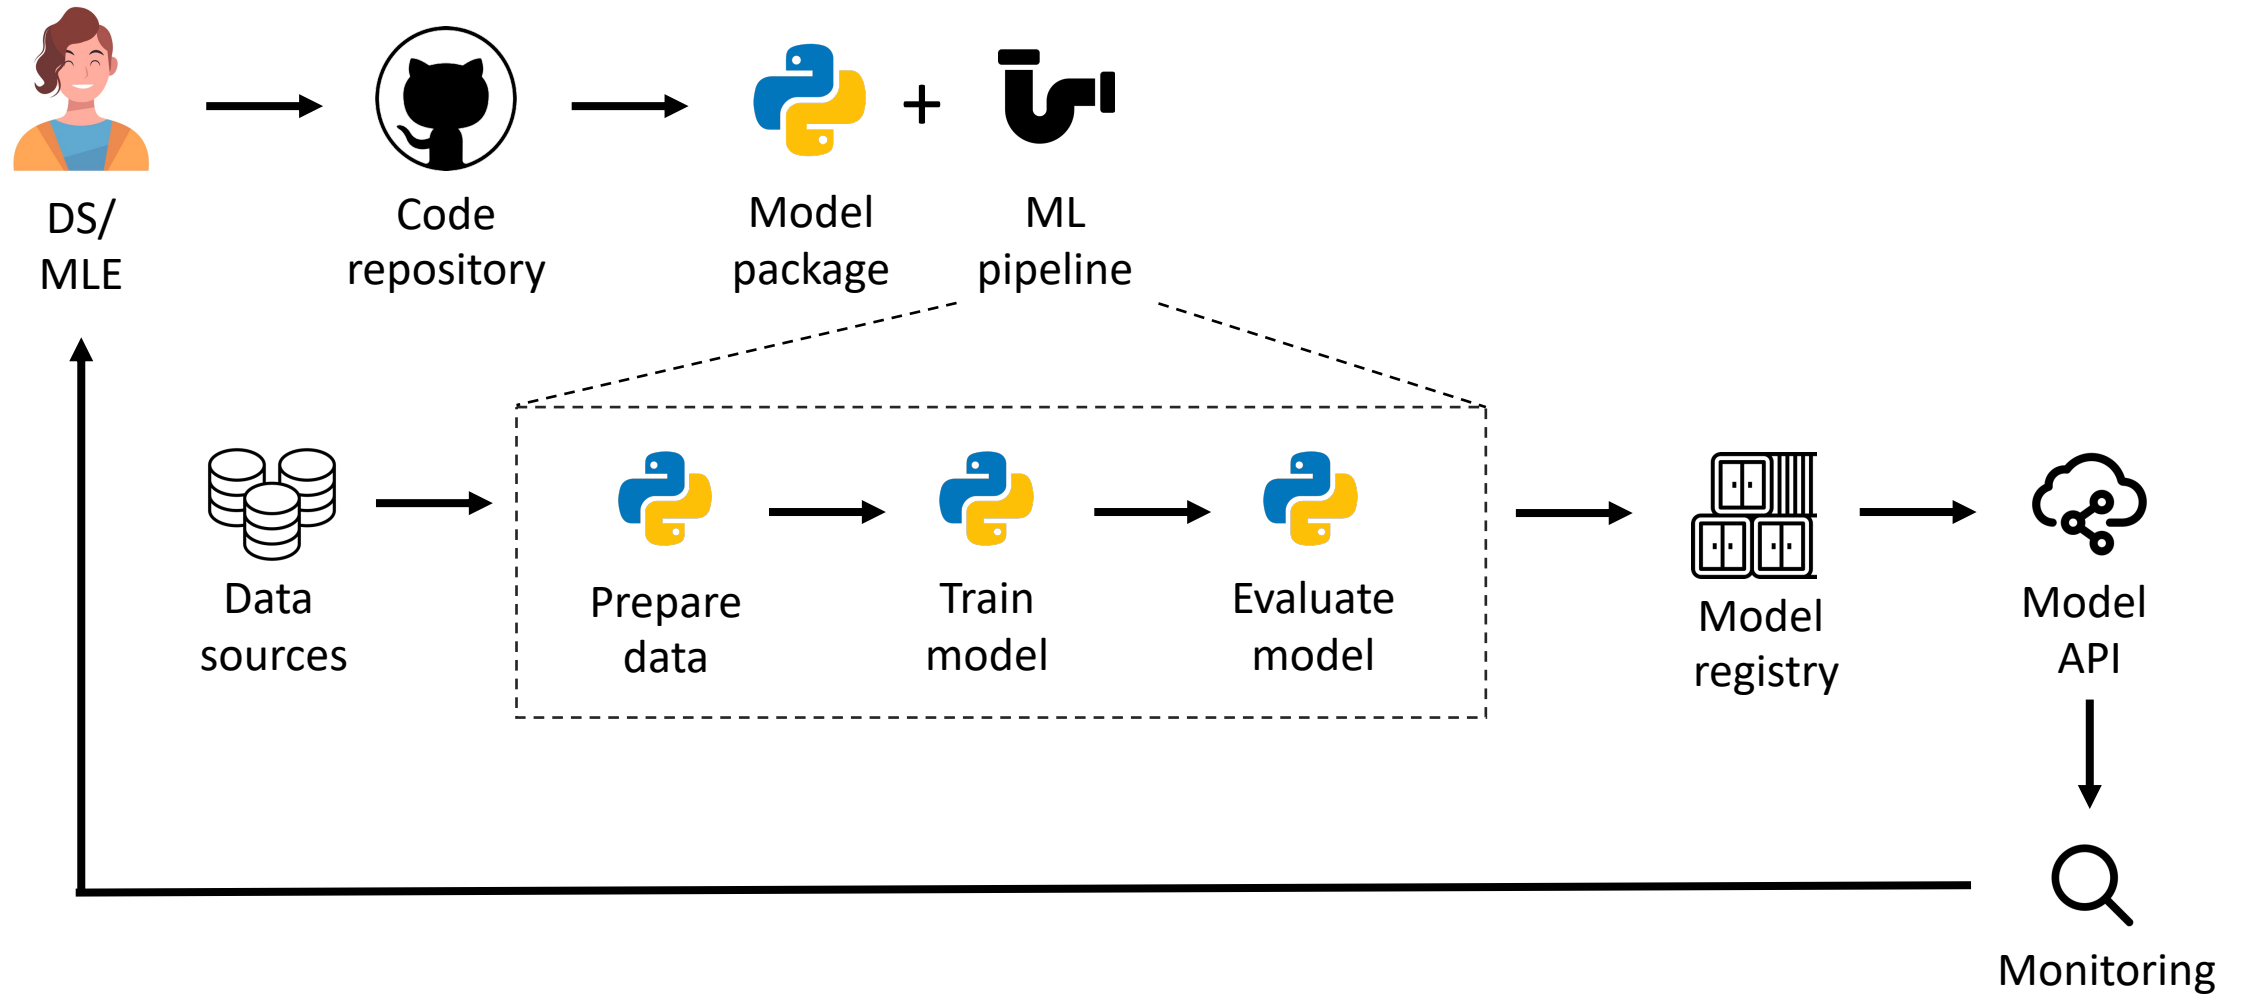
GO   
DATA  
DRIVEN

# What is MLOps?

# What is MLOps?

- Practice that aims to apply lessons from DevOps to build reliable, reproducible and scalable machine learning solutions
- This includes
  - Continuous re-training with re-usable ML pipelines
  - Continuous delivery for new ML models
  - Short feedback loops (monitoring, user feedback)
  - End-to-end ownership of the ML product

# What can this look like in practice?





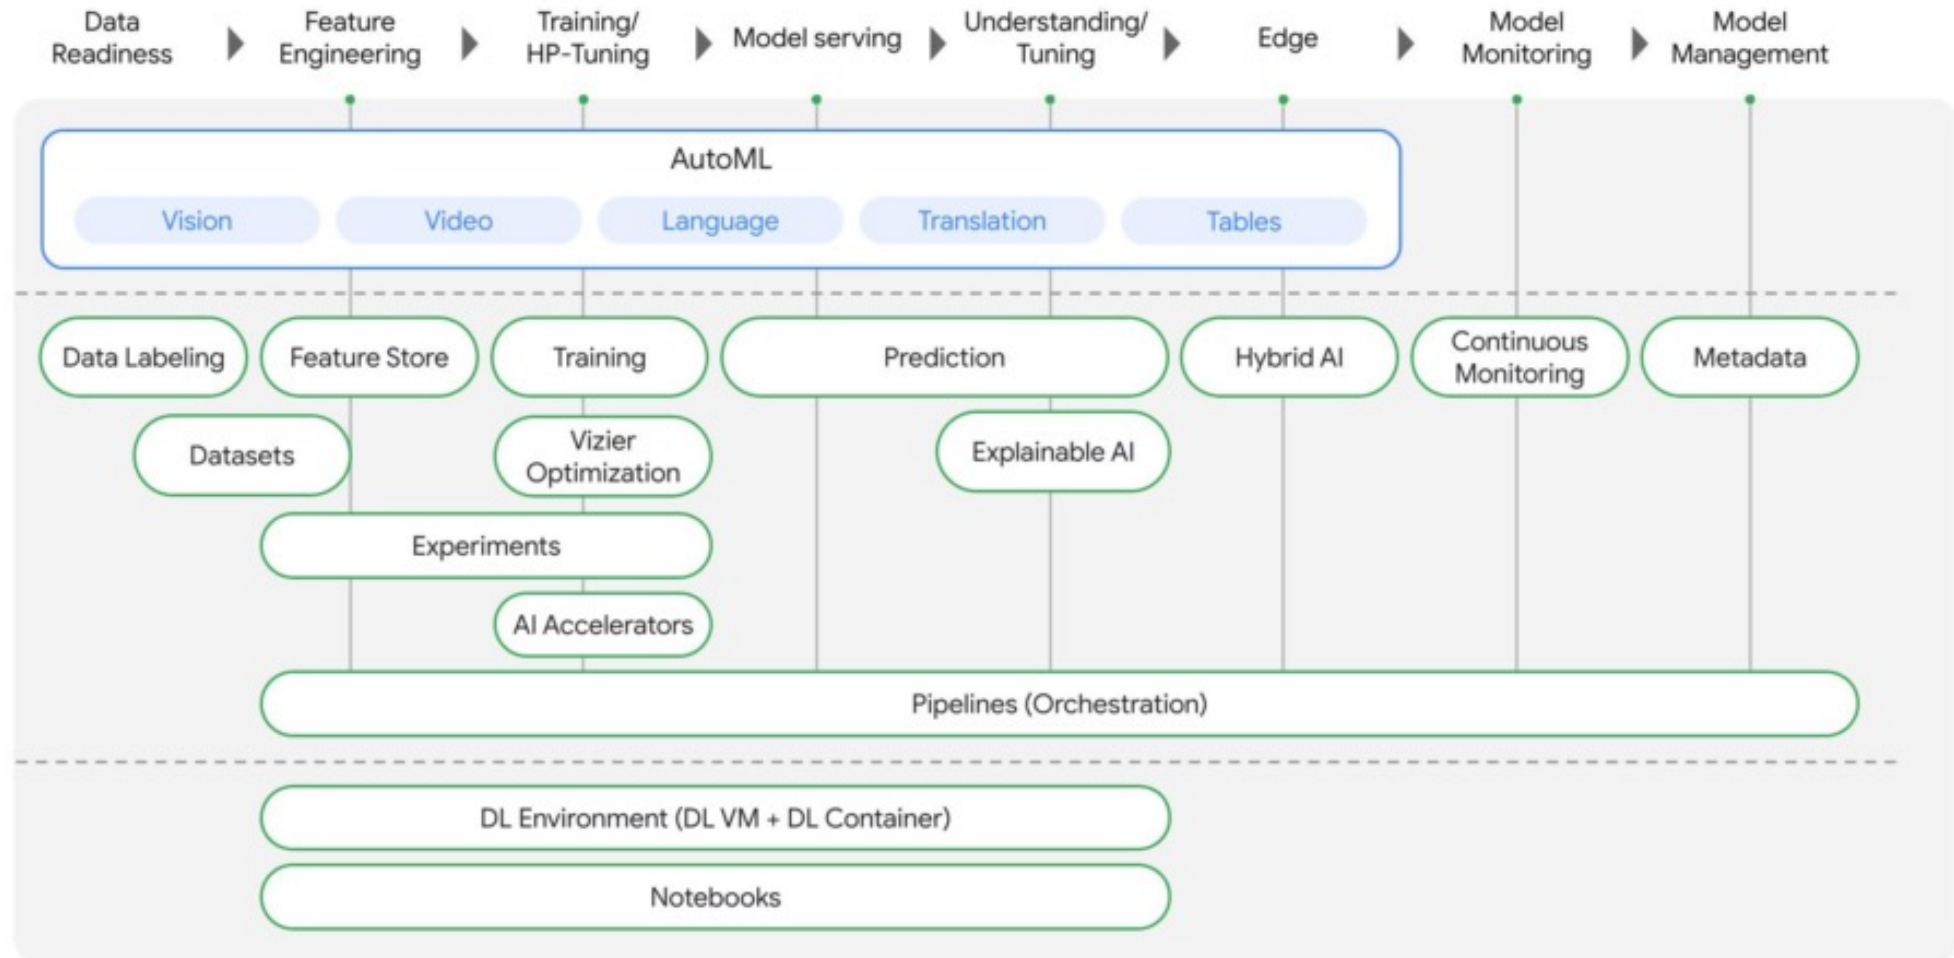
# Vertex AI

GO   
DATA  
DRIVEN

# What is Vertex AI?

- Google's unified platform for ML (formerly *AI Platform*)
- Key features
  - Entire ML workflow (e.g. training, tracking and deploying models) in one unified UI
  - Integrates with open-source frameworks (e.g. Tensorflow) via custom containers
  - Easy integration with other GCP services such as Dataproc, Dataflow, BigQuery, etc.
  - Access to AutoML, pre-trained APIs for video, vision, NLP, etc.

# Vertex AI – Components



# Vertex AI – Demo

Google Cloud Platform

gdd-cb-vertex

Search Products, resources, docs (/)

6

J

Vertex AI

Dashboard

Datasets

Features

Labeling tasks

Workbench

Pipelines

Training

Experiments

Models

Endpoints

Batch predictions

Metadata

Marketplace

Dashboard

Get started with Vertex AI

Vertex AI empowers machine learning developers, data scientists, and data engineers to take their projects from ideation to deployment, quickly and cost-effectively.  
[Learn more](#)

Try an interactive tutorial to learn how to train, evaluate, and deploy a Vertex AI AutoML or custom-trained model  
[VIEW TUTORIALS](#)

Region

europa-west1 (Belgium)

Prepare your training data

Collect and prepare your data, then import it into a dataset to train a model

+ CREATE DATASET

Train your model

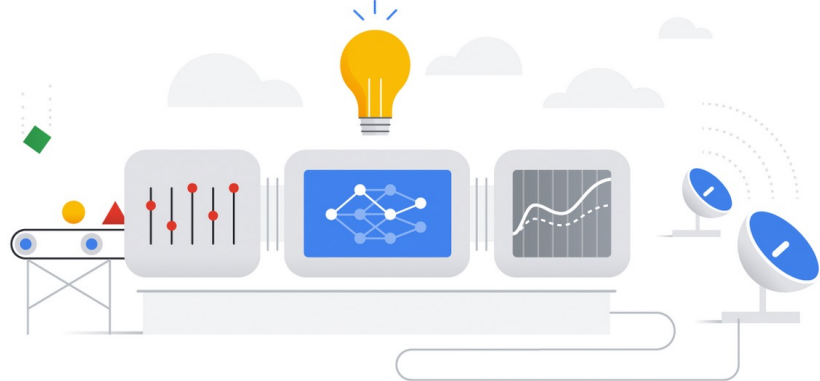
Train a best-in-class machine learning model with your dataset. Use Google's AutoML, or bring your own code.

+ TRAIN NEW MODEL

Get predictions

After you train a model, you can use it to get predictions, either online as an endpoint or through batch requests

+ CREATE BATCH PREDICTION

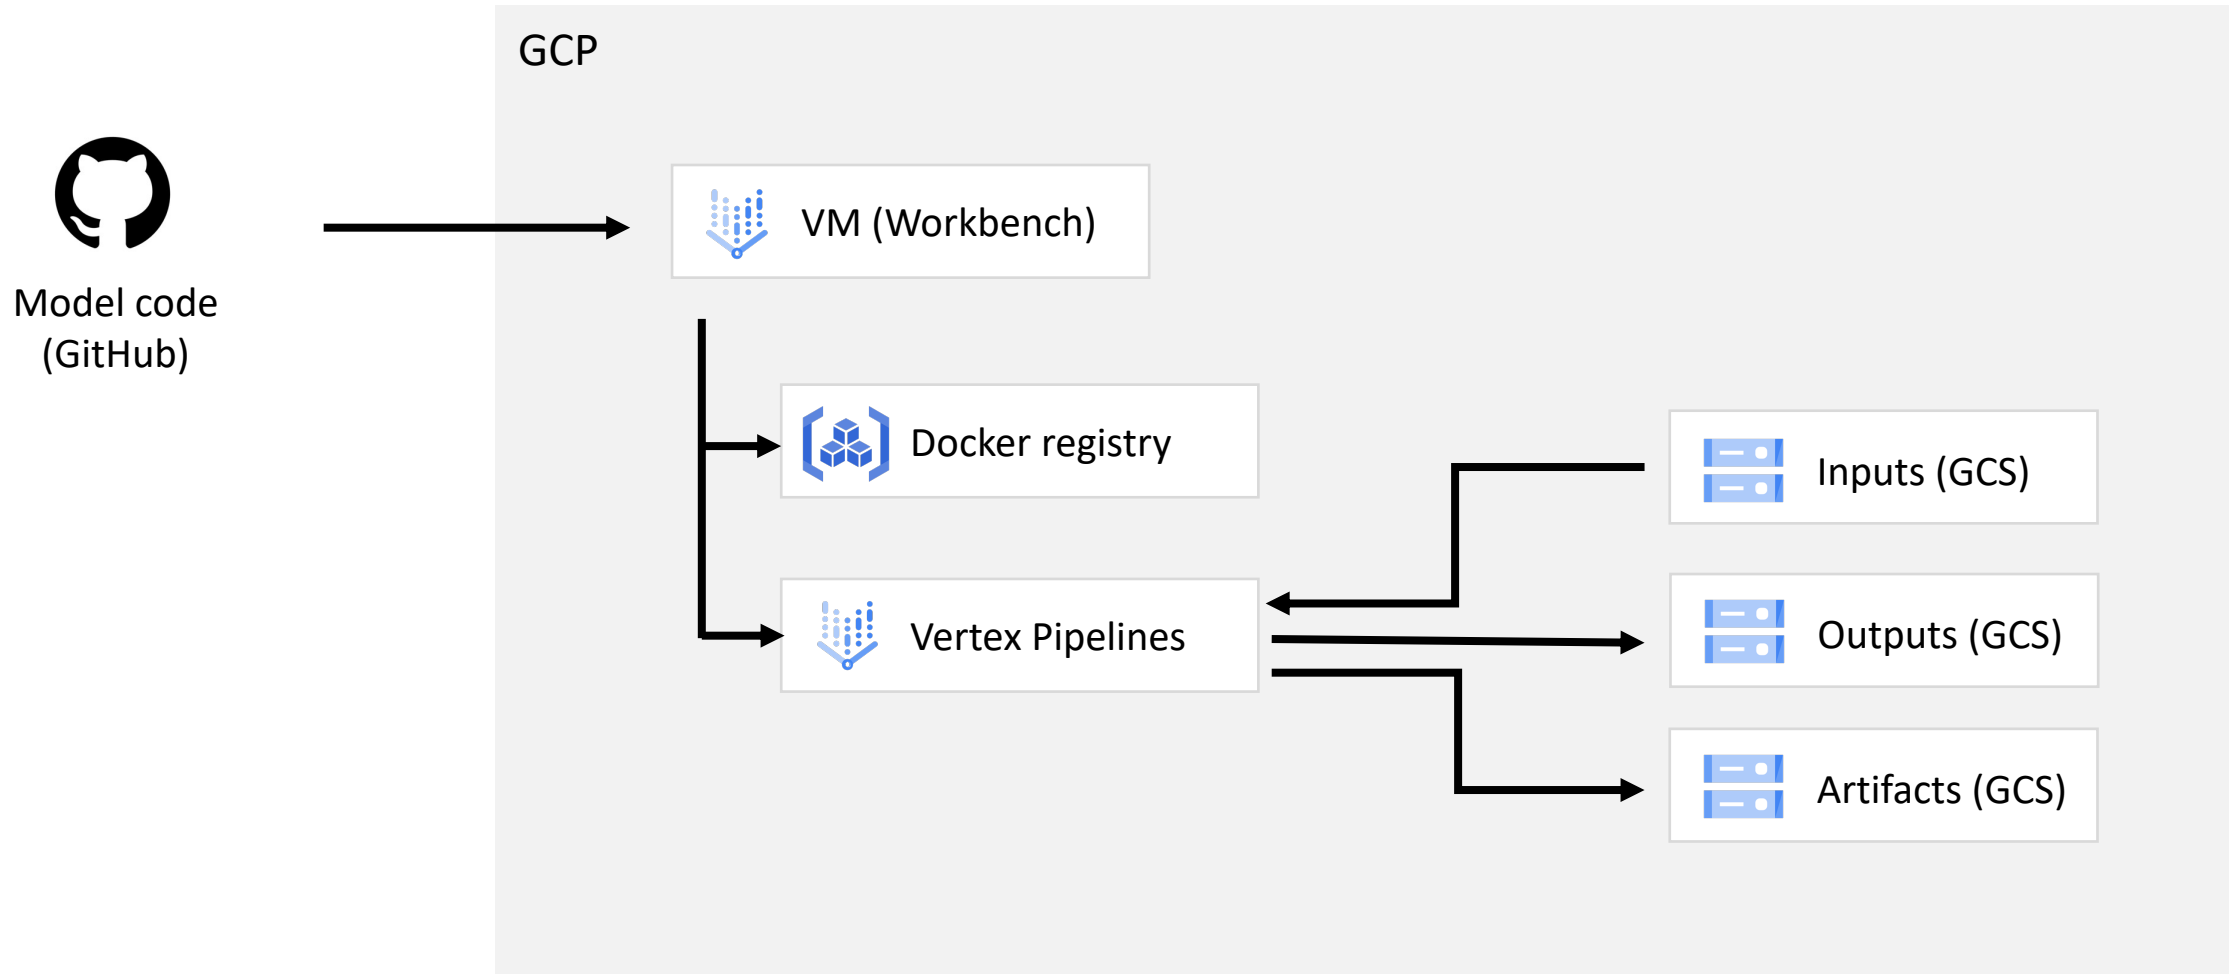


# Getting started

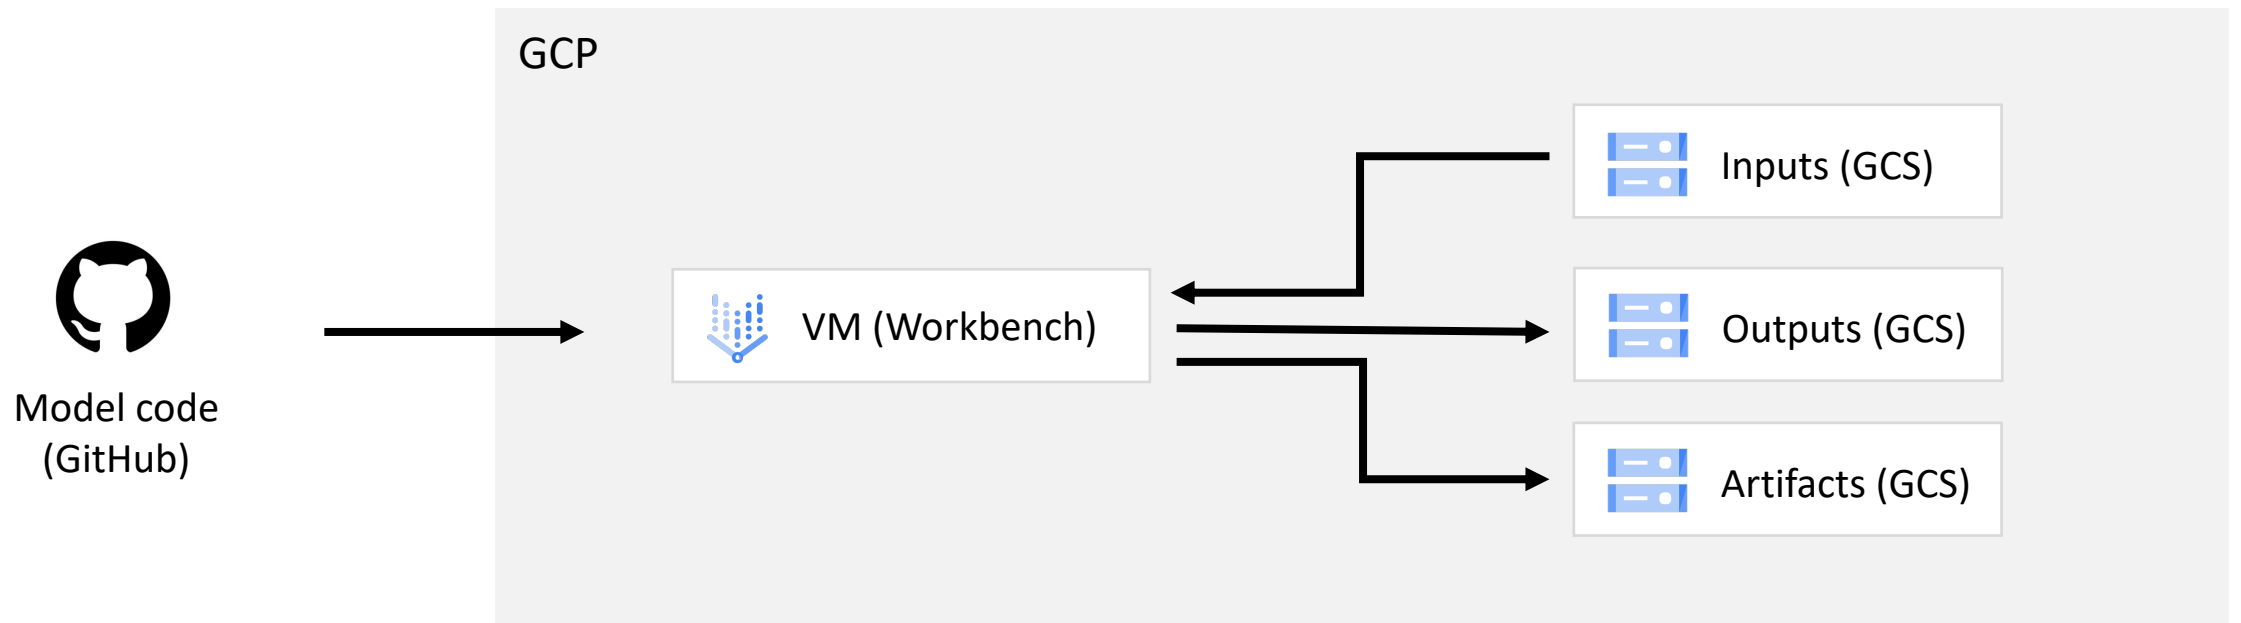
# Use case

- You're working at Fancy Fashion, a sustainable fashion start-up with an app that helps people sell and share second-hand clothing
- A key part of the app is an ML model that automatically analyses uploaded images to automatically assign labels to fashion article
- As a PoC, we're working on a model that takes these images and classifies them into preset categories (e.g. bag, sneaker, etc.)

# Today's target: building a training pipeline



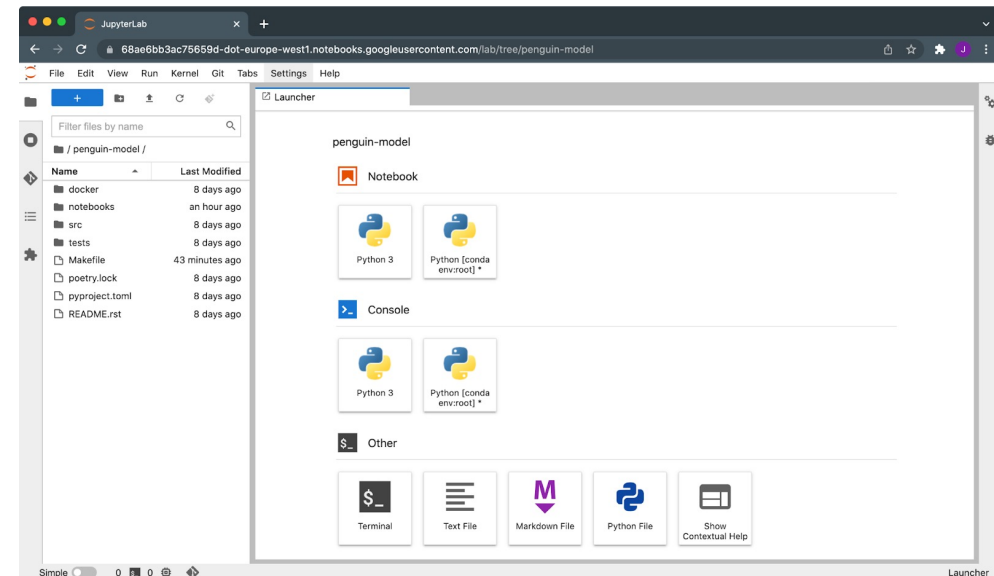
# First: running the model interactively





# What is Vertex Workbench?

- Provides managed VMs with JupyterLab and many common ML frameworks
- Software can be customized with custom VM images or bootstrap script
- Supports different machine sizes + GPUs
- Access to other resources managed transparently using service account
- Accessible over HTTPS (for Jupyter) and SSH (for IDE integration)



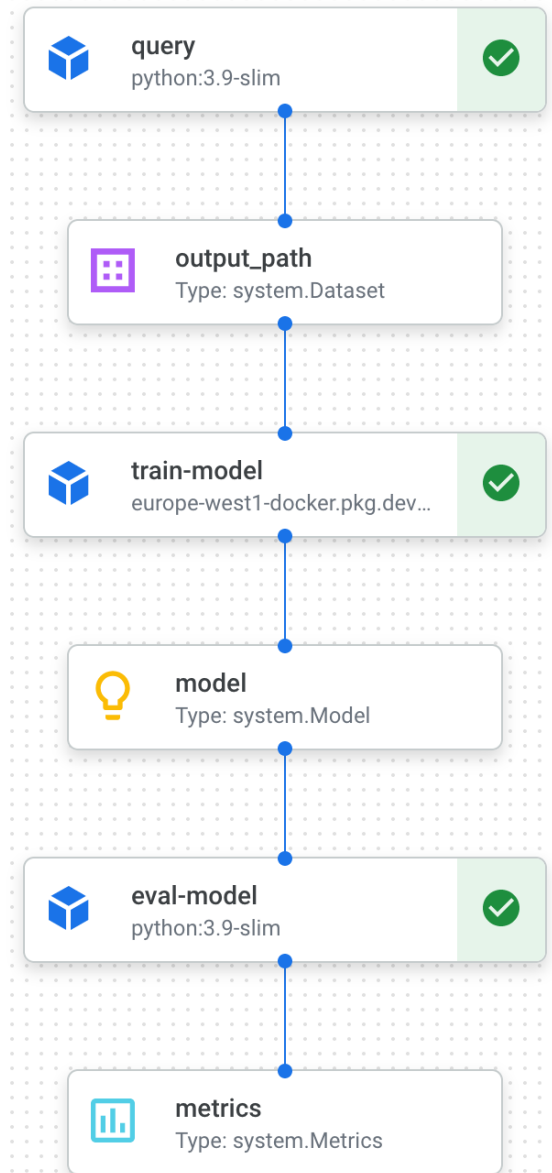
# Exercise: Using Vertex Workbench

- Open the Google Cloud Console in your browser (<https://console.cloud.google.com>)
- Select the 'gdd-cb-vertex' project (top-left)
- Navigate to the 'Vertex AI' section and open the Workbench tab
- Start your VM and open Jupyter Lab
- Clone our repository for the hackathon using git:
  - `git@github.com:jrderuiter/code-breakfast-vertex-ai.git`
- Run through the first exercise (training a model locally)

# Kubeflow pipelines

# What is Kubeflow Pipelines?

- Open-source pipeline SDK for building workflows as DAGs of containerized tasks
- Provides tracking of produced artifacts in a coupled metadata store
- Can be run in Vertex AI using Vertex Pipelines, metadata stored in the ML metadata store



# Defining tasks as components

```
@component(  
    base_image="python:3.9-slim",  
    packages_to_install=["google-cloud-bigquery", "pandas", "pyarrow"],  
    output_component_file="_artifacts/query.yaml",  
)  
def fetch_bigquery(  
    query: str, output_path: OutputPath("Dataset"), project_id: Optional[str] = None  
) -> None:  
    """Runs a query on BigQuery."""  
  
    from google.cloud import bigquery  
  
    client = bigquery.Client(project=project_id)  
    job = client.query(query)  
  
    df = job.to_dataframe()  
    df.to_parquet(output_path)
```

# Combining tasks into a pipeline

```
@kfp.dsl.pipeline(name="penguin")
def pipeline():
    fetch_task = fetch_bigquery(
        "SELECT * FROM bigquery-public-data.ml_datasets.penguins",
        project_id=GCP_PROJECT_ID,
    )

    train_task = (
        train_model(fetch_task.outputs["output_path"])
        # Docs: https://www.kubeflow.org/docs/distributions/gke/pipelines/enable-gpu-and-tpu/
        .set_gpu_limit(1).add_node_selector_constraint(
            "cloud.google.com/gke-accelerator", "nvidia-tesla-k80"
        )
    )

    eval_model(train_task.outputs["model"])
```

# Compiling and running the pipeline

```
compiler.Compiler().compile(  
    pipeline_func=pipeline,  
    package_path="_artifacts/pipeline.json",  
)
```

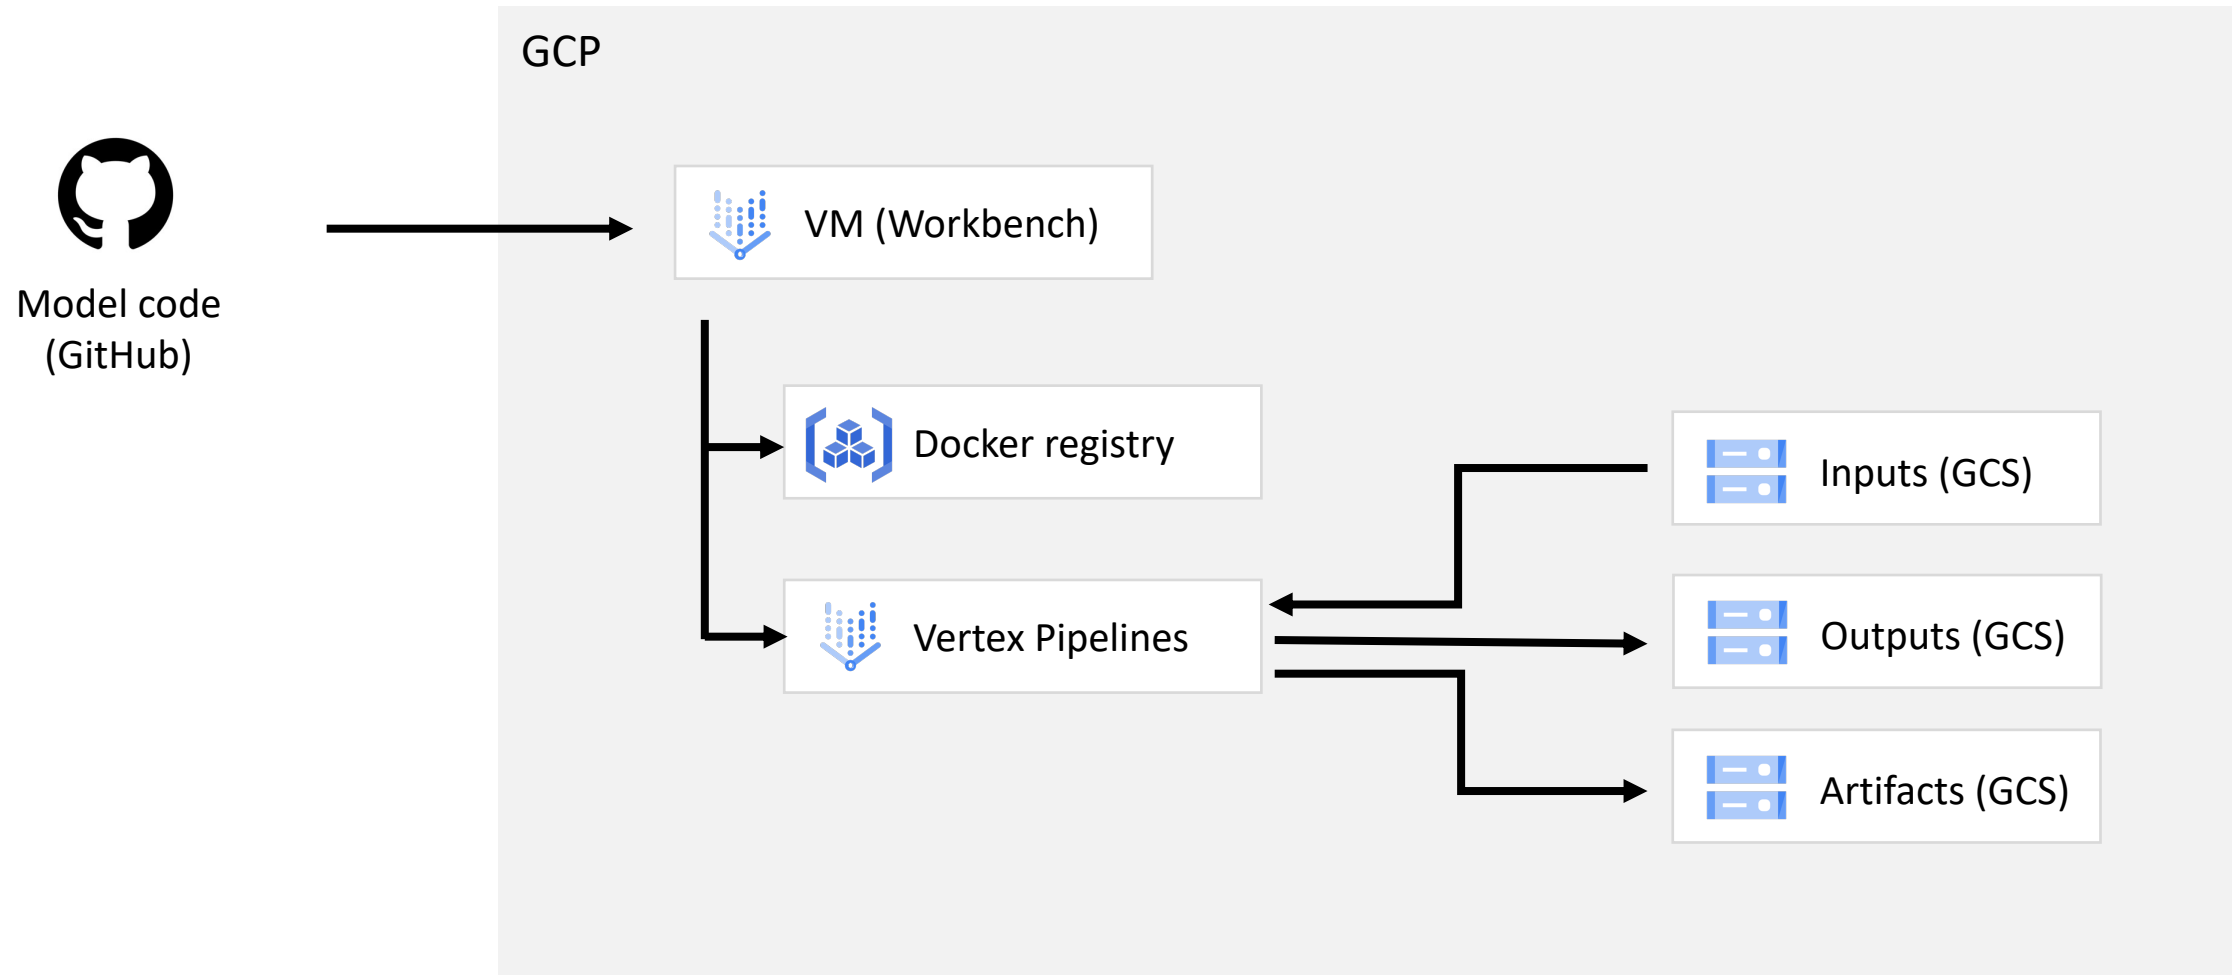
```
from google.cloud.aiplatform.pipeline_jobs import PipelineJob  
  
job = PipelineJob(  
    display_name=f"fancy-fashion-{USER_NAME}",  
    enable_caching=False,  
    template_path="_artifacts/pipeline.json",  
    parameter_values={  
        "train_path": "gs://gdd-cb-vertex-fashion-inputs/train"  
    },  
    pipeline_root=f"gs://gdd-cb-vertex-fashion-artifacts/pipelines",  
    location=GCP_REGION,  
)  
  
job.run(  
    service_account=f"vmd-fashion@gdd-cb-vertex.iam.gserviceaccount.com"  
)
```

# Hackathon: Building a Kubeflow pipeline

- Open the second notebook (2-run-pipeline.ipynb)
- Build and push a Docker image for the package
- Run the provided pipeline and view the results
- Extend the pipeline to include evaluation + prediction steps



# Hackathon: Building an ML pipeline



Questions?